

INF 354

Python

Data Packages

Lecture 11
9 May 2019



The contents of this lecture are as follows:

- Python packages
- A brief introduction to Data Science
- Numpy
- Matplotlib
- Pandas

Packages

- Functions and Methods available in the Python ecosystem can be very useful
- Since there are thousands of functions and methods available, you cannot keep and maintain all those in your project.
- That is where **packages** come in
- Package = Directory of Python Scripts
- Each script in a package is called a module
- These modules specify functions, methods and new Python types (i.e. classes) aimed at solving particular problems
- These packages aren't installed with Python by default – you have to import them
- You can use pip – a package manager for Python
- pip would have been installed when you installed Python

Installing and using packages

- Suppose we want to install the Numpy package. (We'll discuss Numpy in a bit)

- Use pip:

```
SET HTTPS_PROXY=http://[username:password@]proxyserver:portnumber
```

← Configure Proxy

```
pip3 install numpy
```

- Now that Numpy is installed, you can use it in your Python script
 1. Import whole package or specific module from package
 2. Use functions, methods or types provided by package

```
import numpy
```

```
numpy.array([1, 2, 3])
```

```
import numpy as np
```

```
np.array([1, 2, 3])
```

```
from numpy import array
```

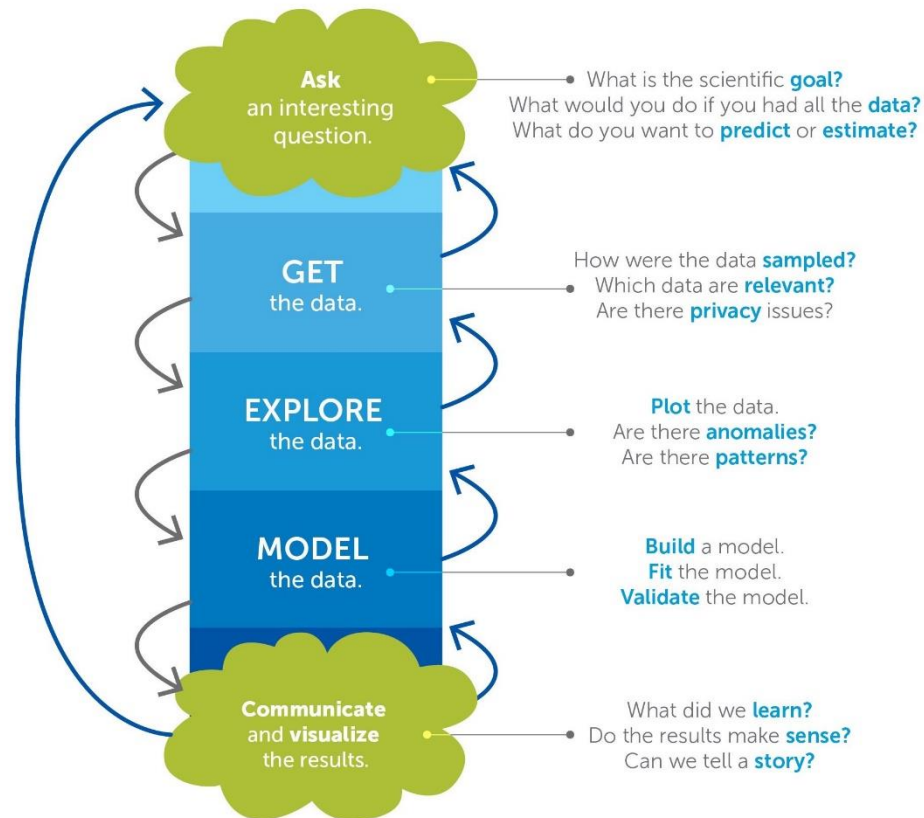
```
array([1, 2, 3])
```

A brief introduction to Data Science

- Data science is the field of study that combines domain expertise, programming skills, and knowledge of math and statistics to extract meaningful insights from data.
- Analysts and business users can then translate these insights into tangible business value.
- It may involve the application of machine learning algorithms to numbers, text, images, video, audio, and more to perform tasks which ordinarily require human intelligence, such as
 - Automated data entry
 - Spam detection
 - Recommendation engines
 - Medical diagnoses
 - Intelligent marketing (customer segmentation, churn prediction, customer value management)
 - Financial analysis (e.g. algorithmic trading, portfolio management, fraud detection, etc.)
 - Predictive maintenance
 - Image recognition

Typical Data Science Process

The Data Science Process



Derived from the work of Joe Blitzstein and Hanspeter Pfister, originally created for the Harvard data science course <http://cs109.org/>.

Limits of built-in Python lists

- Python lists are quite powerful:
 - It can hold collections of different types (string, number, object, list)
 - It can be manipulated with ease (change, add, or remove elements)
- However, data scientists often need to carry out mathematical/statistical operations over entire collections of values – and they want it done **fast**.
- Suppose you have two lists, one holding the heights and one holding the weights of each person in a group. Now, what happens when we try to calculate each person's BMI:

```
In [1]: height = [1.73, 1.68, 1.71, 1.89, 1.79]
In [2]: height
Out[2]: [1.73, 1.68, 1.71, 1.89, 1.79]
In [3]: weight = [65.4, 59.2, 63.6, 88.4, 68.7]
In [4]: weight
Out[4]: [65.4, 59.2, 63.6, 88.4, 68.7]
In [5]: weight / height ** 2
TypeError: unsupported operand type(s) for **: 'list' and 'int'
```

Python does not support
math on entire lists

Looping through the list
will be slow

Numpy

- Numpy (Python package) is short for Numeric Python
- Alternative to the regular Python list: **Numpy Array**
- Numpy array is similar to Python list
- Calculations over entire arrays
- Easy and Fast

```
import numpy as np

np_height = np.array(height)

np_height
array([ 1.73,  1.68,  1.71,  1.89,  1.79])

np_weight = np.array(weight)

np_weight
array([ 65.4,  59.2,  63.6,  88.4,  68.7])
```

```
bmi = np_weight / np_height ** 2

bmi
array([ 21.852,  20.975,  21.75 ,  24.747,  21.441])
```


Unique Numpy Behaviour

- Numpy arrays contain only one type.
- Some operations behave differently to standard Python lists:

```
python_list = [1, 2, 3]
```

```
numpy_array = np.array([1, 2, 3])
```

```
python_list + python_list  
[1, 2, 3, 1, 2, 3]
```

+ operator

Python Lists: Appends two lists

```
numpy_array + numpy_array  
array([2, 4, 6])
```

Numpy Arrays: Adds elements of same index

Quick Numpy Array Queries

Getting value at index of array
(same as Python list)

```
bmi  
array([ 21.852,  20.975,  21.75 ,  24.747,  21.441])  
  
bmi[1]  
20.975
```

Testing all values for some
specific condition

```
bmi > 23  
array([False, False, False,  True, False], dtype=bool)
```

Getting a subset of array
containing only values that meet
some condition

```
bmi[bmi > 23]  
array([ 24.747])
```

2D Numpy Arrays

```
np_2d = np.array([[1.73, 1.68, 1.71, 1.89, 1.79],
                  [65.4, 59.2, 63.6, 88.4, 68.7]])
```

```
np_2d.shape  
(2, 5)
```

2 rows, 5 columns

	0	1	2	3	4	
array([[1.73,	1.68,	1.71,	1.89,	1.79],	0
[65.4,	59.2,	63.6,	88.4,	68.7]])	1

np_2d[0][2]
1.71

np_2d[0,2]
1.71

np_2d[:,1:3]

array([[1.68, 1.71],
 [59.2 , 63.6]])

np_2d[1,:]

array([65.4, 59.2, 63.6, 88.4, 68.7])

Basic Statistics with Numpy

- First step in analysing data is exploration (getting to know your data)
- The list and array examples we've looked at so far have been easy to explore simply by looking at it.
- Data Science usually involves processing thousands, millions, or even billions of numbers

```
In [1]: import numpy as np

In [2]: np_city = ... # Implementation left out

In [3]: np_city
Out[3]:
array([[ 1.64,  71.78],
       [ 1.37,  63.35],
       [ 1.6 ,  55.09],
       ...,
       [ 2.04,  74.85],
       [ 2.04,  68.72],
       [ 2.01,  73.57]])
```

- This is a 2D Numpy array with 5000 rows and 2 columns
- Each row represents a person record
- The first column represents that person's height
- The second column represents that person's weight

- Numpy arrays offer basic statistical methods to explore large data sets such as this one.

Basic Statistics with Numpy

```
In [4]: np.mean(np_city[:,0])
```

```
Out[4]: 1.7472
```

```
In [5]: np.median(np_city[:,0])
```

```
Out[5]: 1.75
```

```
In [6]: np.corrcoef(np_city[:,0], np_city[:,1])
```

```
Out[6]:
```

```
array([[ 1.          , -0.01802],  
       [-0.01803,  1.          ]])
```

```
In [7]: np.std(np_city[:,0])
```

```
Out[7]: 0.1992
```

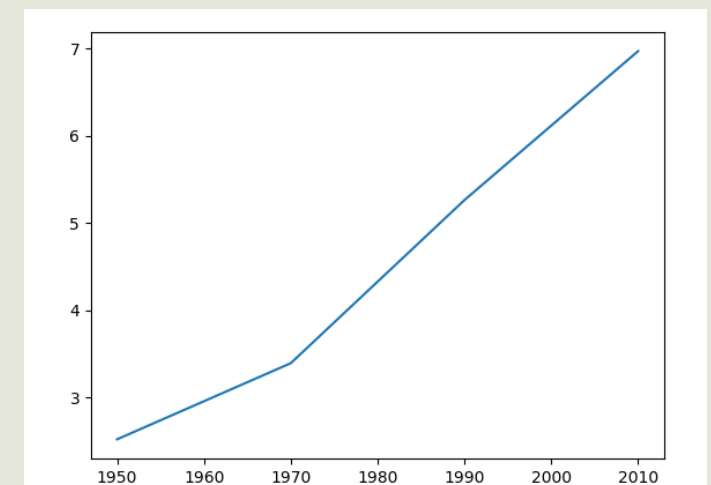
- Mean = arithmetic average of a collection's values
- Median = value lying at the midpoint of a sorted collection.
- Correlation coefficients = measure of the strength of the relationship between two variables. (e.g. height and weight)
- Standard Deviation = measure of how spread out numbers are.
- Numpy arrays also have faster implementations of `sum()` and `sort()`

Matplotlib

- Data visualisation helps with the data analysis process in 2 ways:
 1. Exploring data
 2. Reporting insights
- Matplotlib is a Python package for data visualisation.
- It has set of charting functions, like plot which generates a line chart

```
pip3 install matplotlib
```

```
1 import matplotlib.pyplot as plt
2
3 year = [1950, 1970, 1990, 2010]
4 pop = [2.519, 3.392, 5.263, 6.972]
5
6 plt.plot(year, pop)
7
8 plt.show()
```



Matplotlib - Customisation

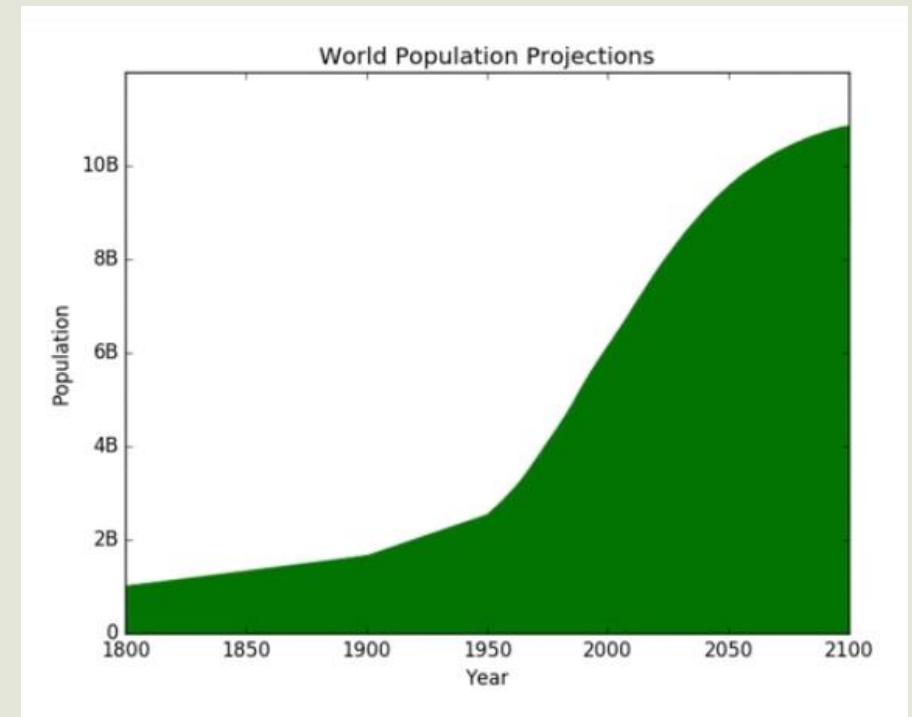
- You can also customise your visualisation by setting certain properties.

```
plt.fill_between(year, population, 0, color='green')

plt.xlabel('Year')
plt.ylabel('Population')
plt.title('World Population Projections')

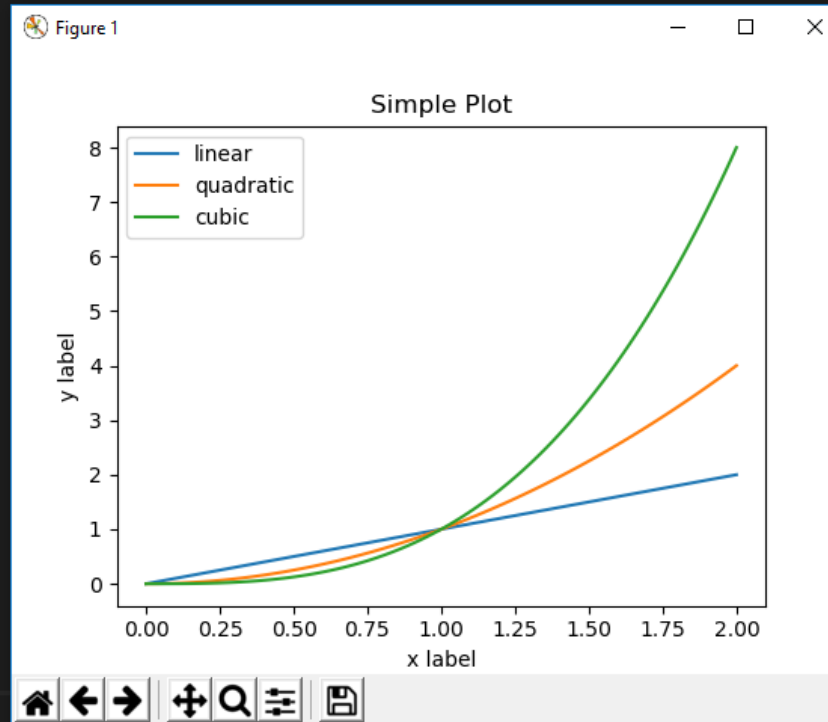
plt.yticks([0, 2, 4, 6, 8, 10],
           ['0', '2B', '4B', '6B', '8B', '10B'])

plt.show()
```



Matplotlib

```
1 import matplotlib.pyplot as plt
2 import numpy as np
3
4 x = np.linspace(0, 2, 100)
5
6 plt.plot(x, x, label='linear')
7 plt.plot(x, x**2, label='quadratic')
8 plt.plot(x, x**3, label='cubic')
9
10 plt.xlabel('x label')
11 plt.ylabel('y label')
12
13 plt.title("Simple Plot")
14
15 plt.legend()
16
17 plt.show()
18
```



For more on matplotlib: <https://matplotlib.org/tutorials/introductory/pyplot.html#sphx-glr-tutorials-introductory-pyplot-py>

Pandas

- Numpy arrays are useful for data manipulation and fast calculations **BUT** they only hold values of a single type
- Big Data is however known for its volume and **variety**, so you cannot expect all data to be of the same type.
- Pandas is a high-level data manipulation package that extends Numpy – it has most of the same easy and fast methods of Numpy but has added methods and works with different data types
- Pandas stores data in a **dataframe** (essentially 2D tables):

```
In [1]: brics = ... # declaration left out
```

```
In [2]: brics
```

```
Out[2]:
```

	country	population	area	capital
BR	Brazil	200	8515767	Brasilia
RU	Russia	144	17098242	Moscow
IN	India	1252	3287590	New Delhi
CH	China	1357	9596961	Beijing
SA	South Africa	55	1221037	Pretoria

Pandas

- Dataframes are typically not built manually. Instead, you import data from an external file (such as CSV file)

```
brics.csv  
  
,country,population,area,capital  
BR,Brazil,200,8515767,Brasilia  
RU,Russia,144,17098242,Moscow  
IN,India,1252,3287590,New Delhi  
CH,China,1357,9596961,Beijing  
SA,South Africa,55,1221037,Pretoria
```



```
brics = pd.read_csv("path/to/brics.csv", index_col = 0)
```



	country	population	area	capital
BR	Brazil	200	8515767	Brasilia
RU	Russia	144	17098242	Moscow
IN	India	1252	3287590	New Delhi
CH	China	1357	9596961	Beijing
SA	South Africa	55	1221037	Pretoria

Pandas

- Data can be accessed in various ways

Column Access

```
In [8]: brics["country"]  
Out[8]:
```

```
BR      Brazil  
RU      Russia  
IN      India  
CH      China  
SA      South Africa  
Name: country, dtype: object
```

```
In [9]: brics.country  
Out[9]:
```

```
BR      Brazil  
RU      Russia  
IN      India  
CH      China  
SA      South Africa  
Name: country, dtype: object
```

```
In [10]: brics["on_earth"] = [True, True, True, True, True]
```

```
In [11]: brics  
Out[11]:
```

	country	population	area	capital	on_earth
BR	Brazil	200	8515767	Brasilia	True
RU	Russia	144	17098242	Moscow	True
IN	India	1252	3287590	New Delhi	True
CH	China	1357	9596961	Beijing	True
SA	South Africa	55	1221037	Pretoria	True

← Add New Column

Row Access

```
In [14]: brics.loc["BR"]  
Out[14]:
```

```
country      Brazil  
population    200  
area         8515767  
capital      Brasilia  
density      23.48585  
on earth     True  
Name: BR, dtype: object
```

Single Cell Access

```
brics.loc["CH", "capital"]  
Beijing
```

```
brics["capital"].loc["CH"]  
Beijing
```

```
brics.loc["CH"]["capital"]  
Beijing
```

```
In [12]: brics["density"] = brics["population"] / brics["area"] * 1000000
```

```
In [13]: brics  
Out[13]:
```

	country	population	area	capital	on_earth	density
BR	Brazil	200	8515767	Brasilia	True	23.485847
RU	Russia	144	17098242	Moscow	True	8.421918
IN	India	1252	3287590	New Delhi	True	380.826076
CH	China	1357	9596961	Beijing	True	141.398928
SA	South Africa	55	1221037	Pretoria	True	45.043680

← Add New Calculated Column

Video

- <https://www.youtube.com/watch?v=a9UrKTVeeZA&t=835s>

*Note the producer of this video works in a different IDE (namely Jupyter). However, he uses Python so the code will work just as well in VS Code.